

Attention in Graph Neural Networks

Prof. O-Joun Lee

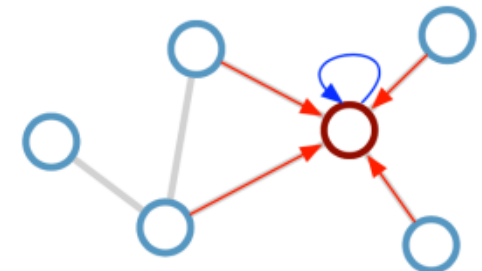
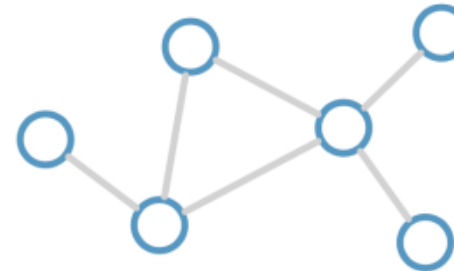
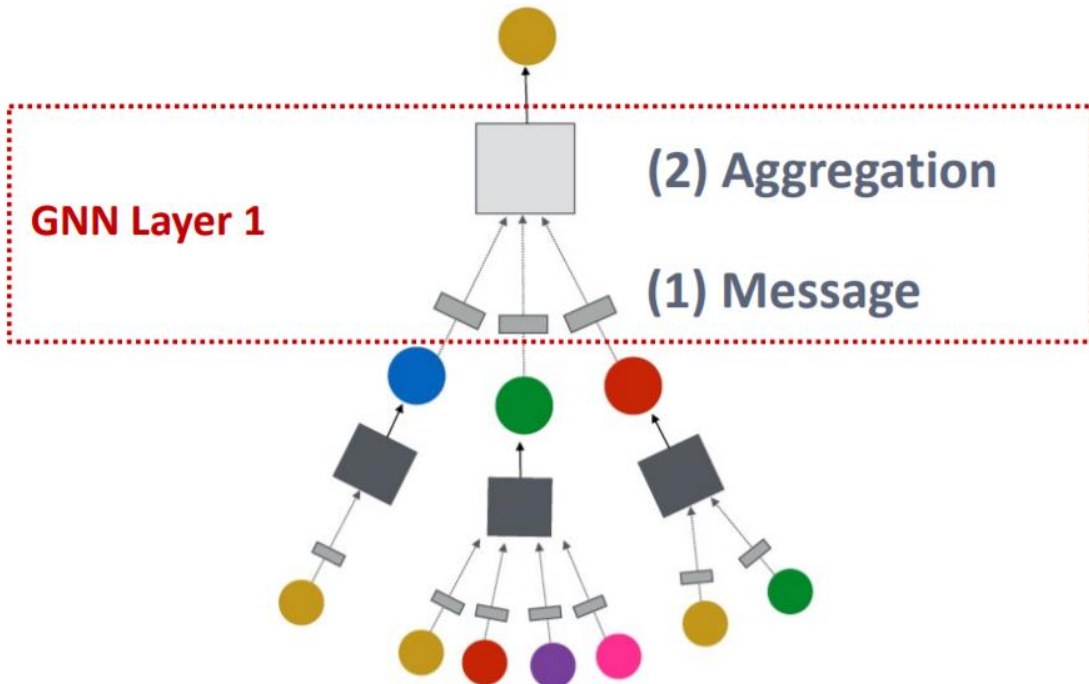
Dept. of Artificial Intelligence,
The Catholic University of Korea
ojlee@catholic.ac.kr

Contents



- Graph neural network issues
- Attention in Graph neural networks
- Attention in Heterogeneous graphs
- GAT sample code


- GNN Layer = **Message** + **Aggregation**
 - Message COMPUTATION
 - how to make each neighborhood node as embedding?
 - Message AGGERGATION
 - how to combine those embeddings?



Update rule:
$$\mathbf{h}_i^{(l+1)} = \sigma \left(\mathbf{h}_i^{(l)} \mathbf{W}_0^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{c_{ij}} \mathbf{h}_j^{(l)} \mathbf{W}_1^{(l)} \right)$$

- **Intuition:** Each node will create a message, which will be sent to other nodes later
- **Example:** A Linear layer $\mathbf{m}_u^{(l)} = \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)}$
 - Multiply node features with weight matrix $\mathbf{W}^{(l)}$

Message function: $\mathbf{m}_u^{(l)} = \text{MSG}^{(l)} \left(\mathbf{h}_u^{(l-1)} \right)$





- **Intuition:** Each node will aggregate the messages from node v 's neighbors

$$\mathbf{h}_v^{(l)} = \text{AGG}^{(l)} \left(\left\{ \mathbf{m}_u^{(l)}, u \in N(v) \right\} \right)$$

- **Example:** Sum(\cdot), Mean(\cdot) or Max(\cdot) aggregator

$$\mathbf{h}_v^{(l)} = \text{Sum}(\{\mathbf{m}_u^{(l)}, u \in N(v)\})$$



- **Issue:** Information from node v itself could get lost
 - Computation of $\mathbf{h}_v^{(l)}$ does not directly depend on $\mathbf{h}_v^{(l-1)}$

- **Solution:** Include $\mathbf{h}_v^{(l-1)}$ when computing $\mathbf{h}_v^{(l)}$

- (1) **Message:** compute message from node v itself

$$\text{●} \text{●} \text{●} \quad \mathbf{m}_u^{(l)} = \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)} \qquad \text{●} \quad \mathbf{m}_v^{(l)} = \mathbf{B}^{(l)} \mathbf{h}_v^{(l-1)}$$

- (2) **Aggregation:** After aggregating from neighbors, we can aggregate the message from node v itself
 - Via concatenation or summation

$$\mathbf{h}_v^{(l)} = \text{CONCAT} \left(\underbrace{\text{AGG} \left(\left\{ \mathbf{m}_u^{(l)}, u \in N(v) \right\} \right)}_{\text{First aggregate from neighbors}}, \underbrace{\mathbf{m}_v^{(l)}}_{\text{Then aggregate from node itself}} \right)$$

- Pure Graph Convolutional Networks (GCN)

$$\mathbf{h}_v^{(l)} = \sigma \left(\sum_{u \in N(v)} \mathbf{W}^{(l)} \frac{\mathbf{h}_u^{(l-1)}}{|N(v)|} \right)$$

- **Message:** Each neighbour u :


$$\mathbf{m}_u^{(l)} = \frac{1}{|N(v)|} \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)}$$

- **Aggregation:** Sum over messages from neighbors, then apply activation

$$\mathbf{h}_v^{(l)} = \sigma \left(\text{Sum} \left(\left\{ \mathbf{m}_u^{(l)}, u \in N(v) \right\} \right) \right)$$


→ All neighbors $u \in N(v)$ are equally important to node v

equally important to v




Not all node's neighbors are equally important

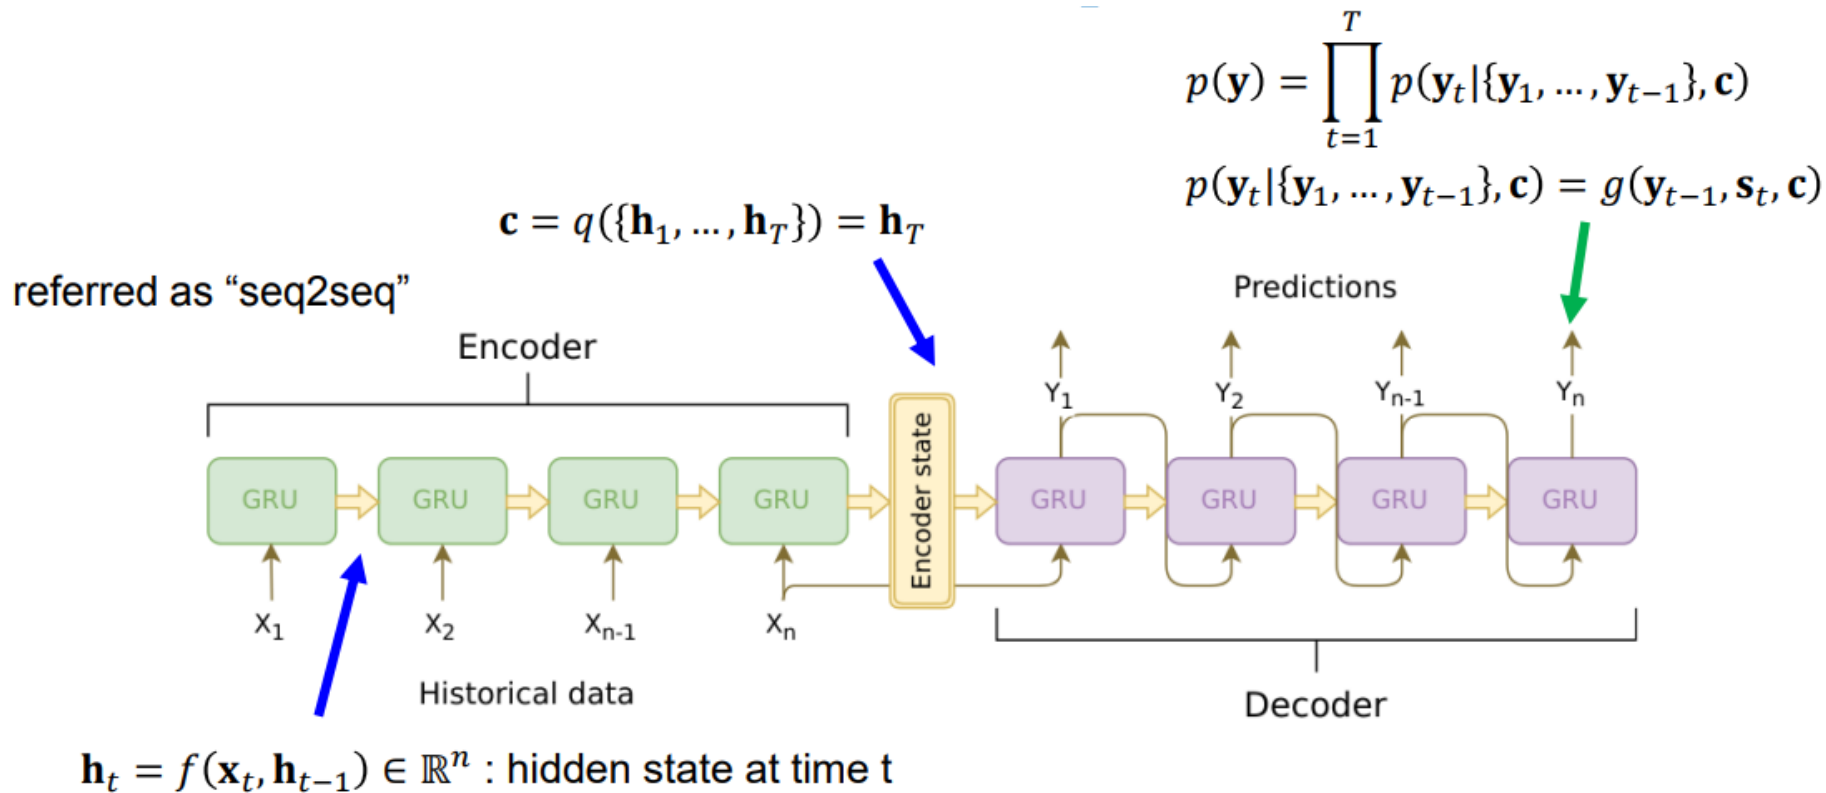
- Attention is a mechanism that allows a network to focus on certain parts of the input when processing it
- The attention focuses on the important parts of the input data and fades out the rest.
 - **Idea:** the NN should devote more computing power on that small but important part of the data.
 - Which part of the data is more important depends on the context and is learned through training.




- We deduce something by paying attention to something that is relatively more important.




- RNN encoder-decoder for neural machine translation:
 - In capability of remembering long sentences : Often it has forgotten the first part once it completes processing the whole input. The attention mechanism was born to resolve this problem.



- GNN compute node representations from representations of neighbours
- Nodes can have largely different neighbourhood sizes
- Not all neighbours have relevant information for a certain node
- Attention mechanism allow to adaptively weight the contribution of each neighbour when updating a node




- **Attention means:** assign an attention coefficient to each neighbor, indicating the importance of that neighbor's features for the feature update of the node.




[Figure from Veličković et al. (ICLR 2018)]

- The key difference between GAT and GCN is how the information from the one-hop neighborhood is aggregated.


Vanilla GCN updates information of neighbor nodes with same importance



Attention mechanism enables GCN to update nodes with different importance.



- In Graph Attention Networks (GATs), the concept of multiple attention heads is similar to the idea of multiple filters in Convolutional Neural Networks (CNNs).
- Each attention head can potentially learn to pay attention to different types of neighborhood information.




- Input node features: Each node in the graph has a feature vector.

$$\mathbf{h} = \{\vec{h}_1, \vec{h}_2, \dots, \vec{h}_N\}, \vec{h}_i \in \mathbb{R}^F$$

- Calculate energy (co-efficiency) between two nodes

$$e_{ij} = a(\mathbf{W}\vec{h}_i, \mathbf{W}\vec{h}_j)$$

a: attention function



- Attention score (over the neighbors): Normalize over all the neighbors


$$\alpha_{ij} = \frac{\exp \left(\text{LeakyReLU} \left(\vec{\mathbf{a}}^T [\mathbf{W} \vec{h}_i \| \mathbf{W} \vec{h}_j] \right) \right)}{\sum_{k \in \mathcal{N}_i} \exp \left(\text{LeakyReLU} \left(\vec{\mathbf{a}}^T [\mathbf{W} \vec{h}_i \| \mathbf{W} \vec{h}_k] \right) \right)}$$

- Multi-head attention
 - Feature concatenation

$$\vec{h}'_i = \parallel_{k=1}^K \sigma \left(\sum_{j \in \mathcal{N}_i} \alpha_{ij}^k \mathbf{W}^k \vec{h}_j \right)$$


- Feature averaging (for the final layer)

$$\vec{h}'_i = \sigma \left(\frac{1}{K} \sum_{k=1}^K \sum_{j \in \mathcal{N}_i} \alpha_{ij}^k \mathbf{W}^k \vec{h}_j \right)$$

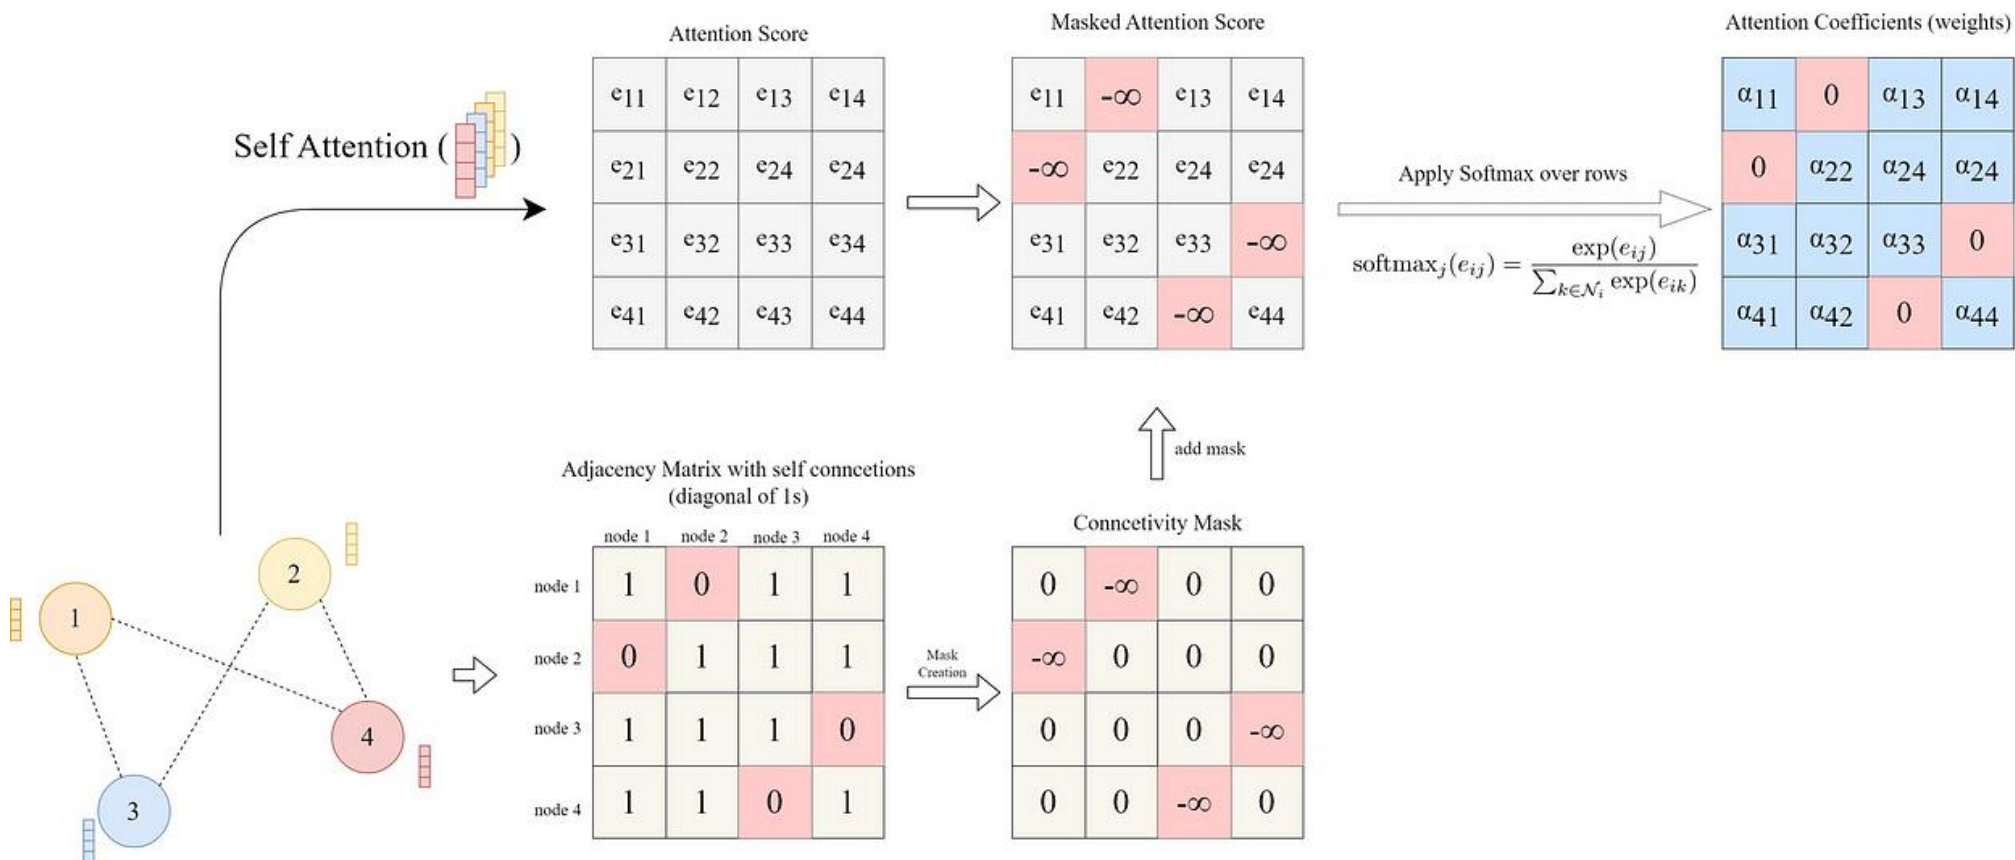


- Pros:
 - No need to score intermediate edge-based activation vectors (when using dot product attention)
 - Slower than GCNs but faster than GNNs with edge embeddings
- Cons:
 - Can be more difficult to optimize


- The whole operation is illustrated below:



- Applying masking mechanism to the masked attention score, then apply Softmax function:

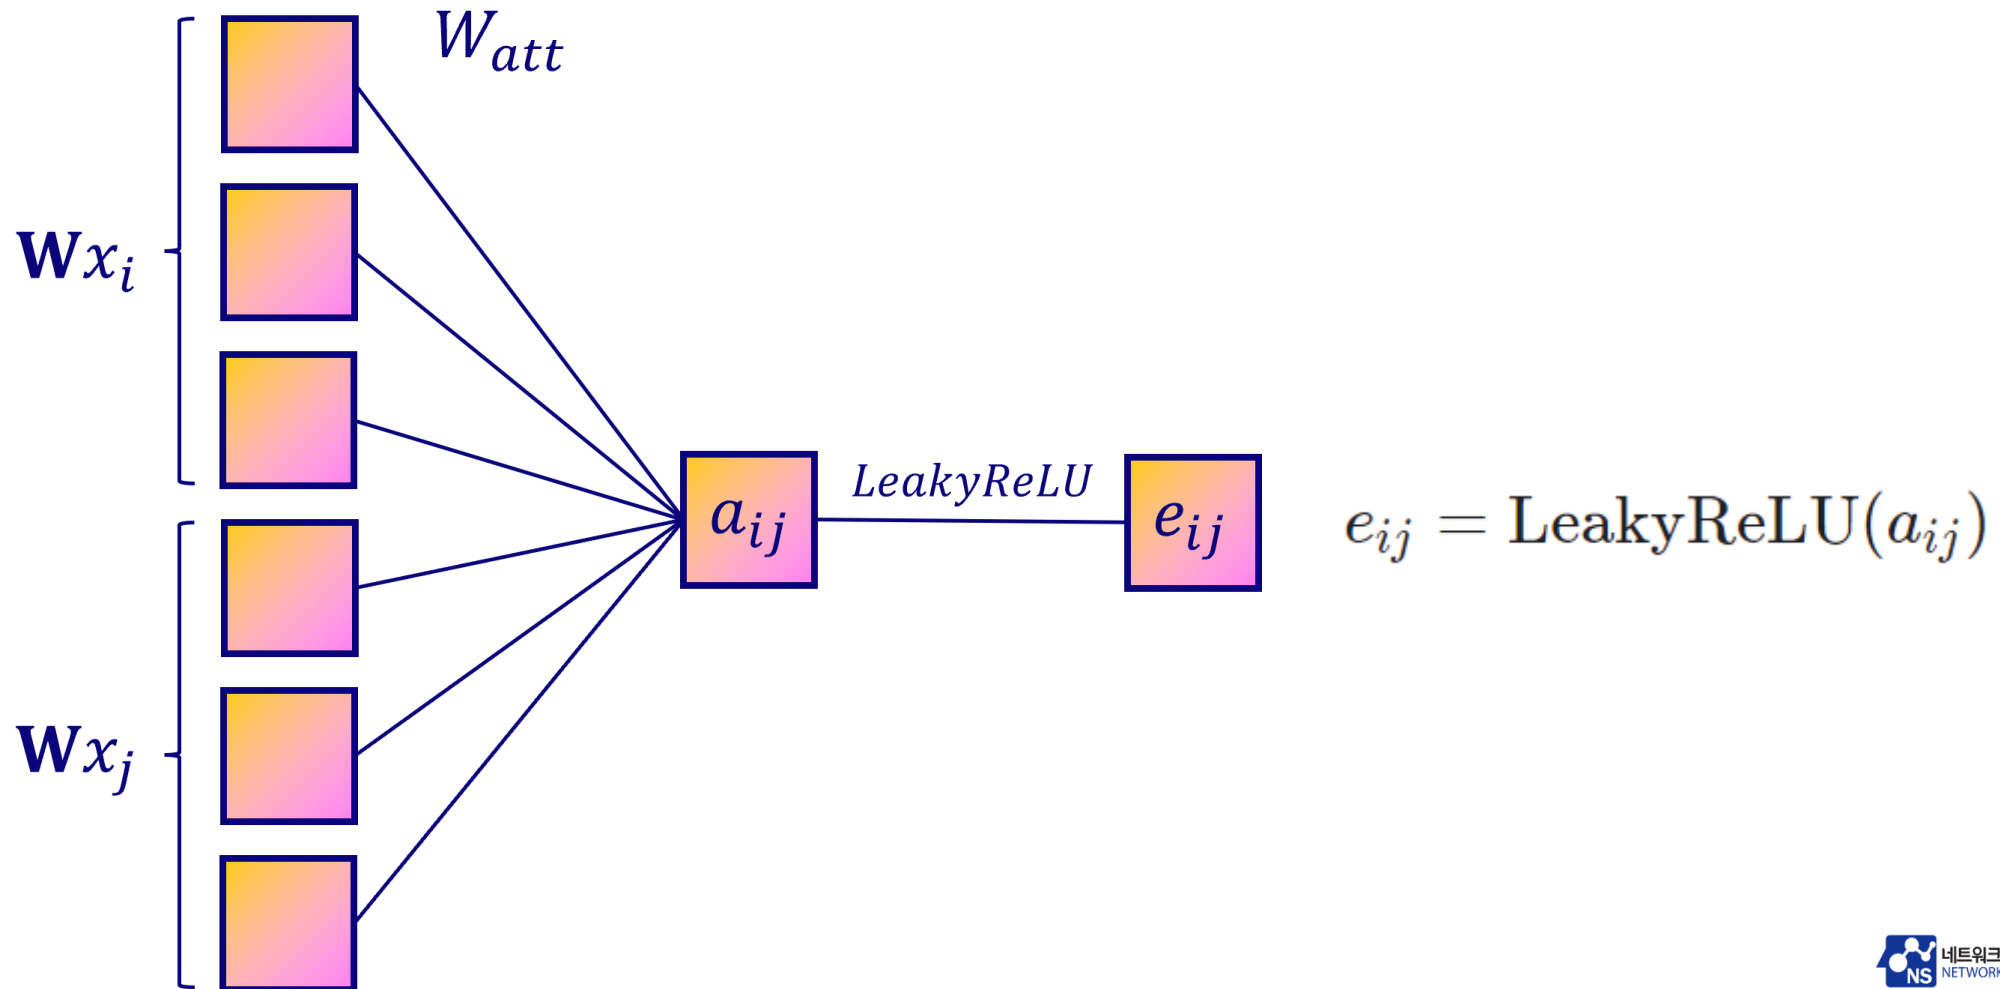


- **Linear transformation:** To calculate the attention coefficient, we need to consider pairs of nodes. An easy way to create these pairs is to concatenate attribute vectors from both nodes.

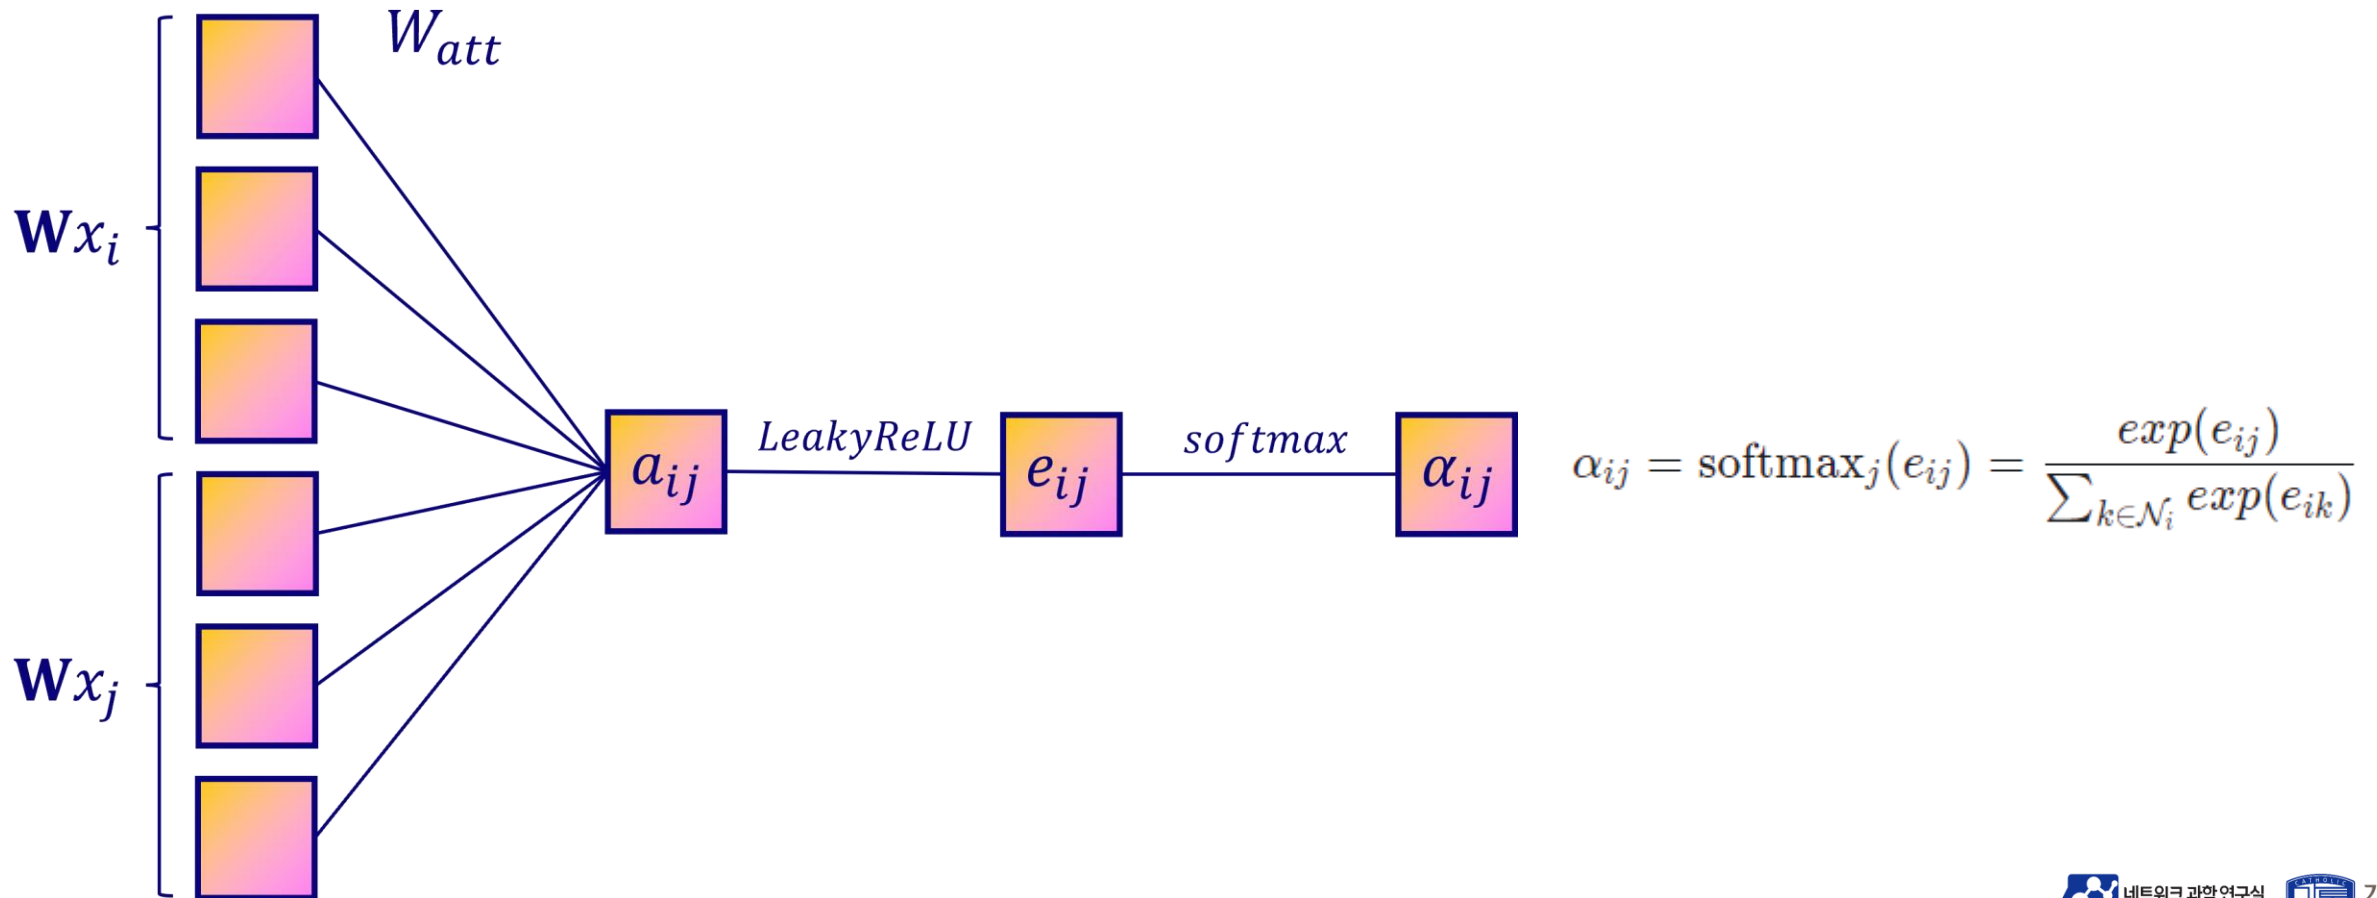


$$a_{ij} = W_{att}^t [\mathbf{W}x_i \parallel \mathbf{W}x_j]$$

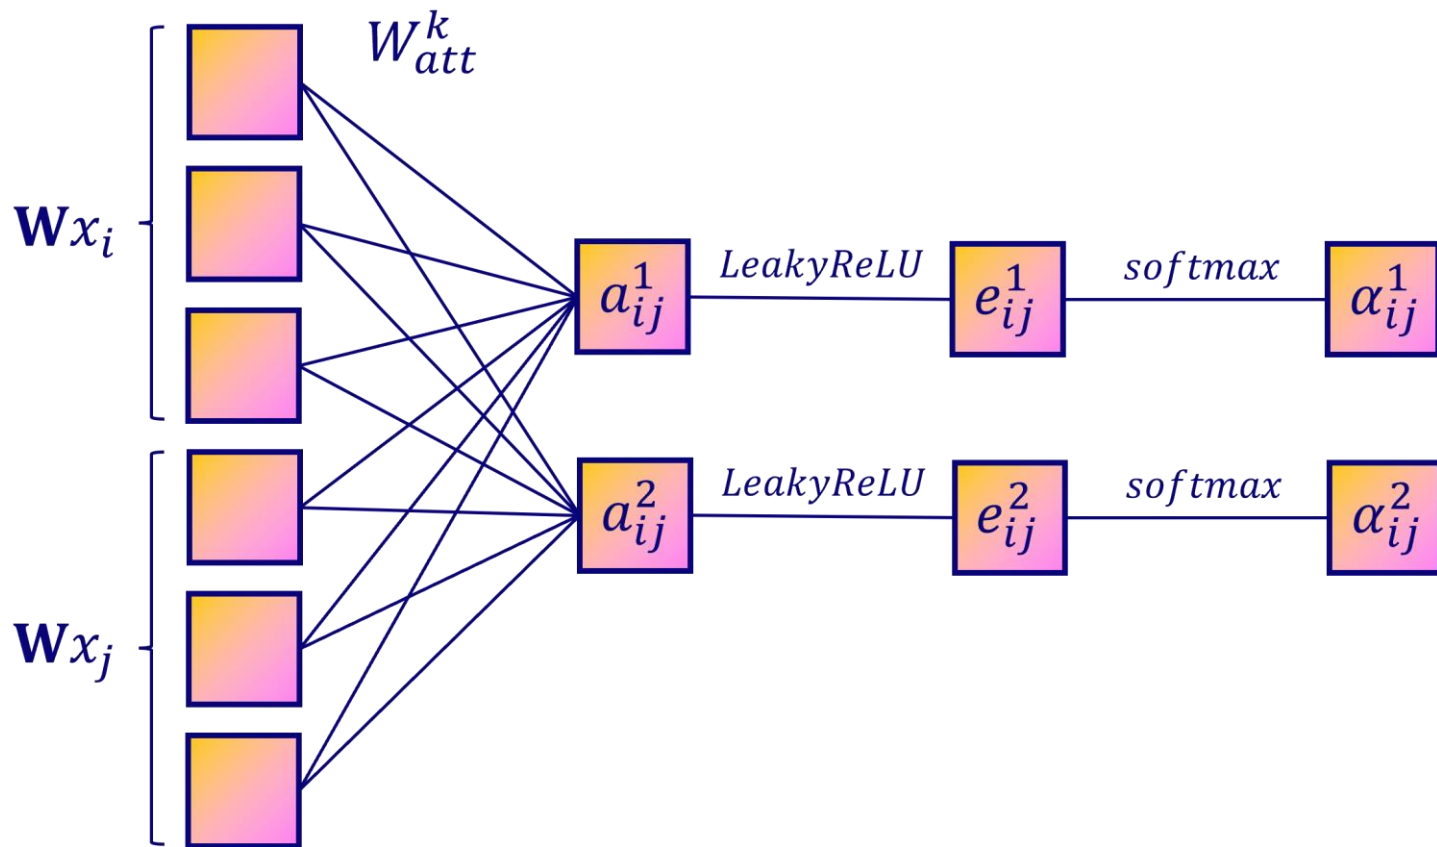
- **Activation function:** add nonlinearity with an activation function. In this case, the paper's authors chose the LeakyReLU LeakyReLU function.



- Softmax normalization: The output of our neural network is not normalized, which is a problem since we want to compare these coefficients.
- A common way to do it with neural networks is to use the softmax function.



- Multi-head attention: In GATs, multi-head attention consists of replicating the same three steps several times in order to average or concatenate the results.



Average:

$$h_i = \frac{1}{n} \sum_{k=1}^n h_i^k$$

Concatenation:

$$h_i = \parallel_{k=1}^n h_i^k$$

[🏠](#) / [torch_geometric.nn](#) / `conv.GATConv`

conv.GATConv

```
class GATConv ( in_channels: Union[int, Tuple[int, int]], out_channels: int, heads: int = 1, concat: bool =
True, negative_slope: float = 0.2, dropout: float = 0.0, add_self_loops: bool = True, edge_dim:
Optional[int] = None, fill_value: Union[float, Tensor, str] = 'mean', bias: bool = True, **kwargs )
    [source]
```

Bases: `MessagePassing`

The graph attentional operator from the “Graph Attention Networks” paper

$$\mathbf{x}'_i = \alpha_{i,i} \Theta \mathbf{x}_i + \sum_{j \in \mathcal{N}(i)} \alpha_{i,j} \Theta \mathbf{x}_j,$$

where the attention coefficients $\alpha_{i,j}$ are computed as

$$\alpha_{i,j} = \frac{\exp(\text{LeakyReLU}(\mathbf{a}^\top [\Theta \mathbf{x}_i \parallel \Theta \mathbf{x}_j]))}{\sum_{k \in \mathcal{N}(i) \cup \{i\}} \exp(\text{LeakyReLU}(\mathbf{a}^\top [\Theta \mathbf{x}_i \parallel \Theta \mathbf{x}_k]))}.$$

```
46  class GAT(torch.nn.Module):
47      def __init__(self, num_features, num_classes, dims, drop=0.0):
48          super(GAT, self).__init__()
49          heads = 8
50          self.conv1 = GATConv(num_features, dims, heads=heads, dropout=0.3, concat=False)
51          # On the Pubmed dataset, use heads=8 in conv2.
52          self.conv2 = GATConv(dims, num_classes, heads=heads, concat=False,
53                               dropout=0.3)
54          self.drop = torch.nn.Dropout(p=drop)
55      def forward(self, x, edge_index):
56          x = F.elu(self.conv1(x, edge_index))
57          x = self.drop(x)
58          x = self.conv2(x, edge_index)
59          return F.log_softmax(x, dim=1), x
```

➤ Let's try some simple GAT code in the sample code file

- GATv2s is similar to GAT.
- The GATv2 operator fixes the static attention problem of the standard GAT.
 - Static attention is when the attention to the key nodes has the same rank (order) for any query node.
 - GAT computes attention from query node i to key node j :



$$\begin{aligned} e_{ij} &= \text{LeakyReLU}\left(\mathbf{a}^\top \left[\mathbf{W} \vec{h}_i \parallel \mathbf{W} \vec{h}_j \right]\right) \\ &= \text{LeakyReLU}\left(\mathbf{a}_1^\top \mathbf{W} \vec{h}_i + \mathbf{a}_2^\top \mathbf{W} \vec{h}_j\right) \end{aligned}$$

GAT



$$\begin{aligned} e_{ij} &= \mathbf{a}^\top \text{LeakyReLU}\left(\mathbf{W} \left[\vec{h}_i \parallel \vec{h}_j \right]\right) \\ &= \mathbf{a}^\top \text{LeakyReLU}\left(\mathbf{W}_l \vec{h}_i + \mathbf{W}_r \vec{h}_j\right) \end{aligned}$$

GATv2

- The GATv2 model performs better than the first version GAT, because it uses a dynamic graph attention variant that has a universal approximator attention function, it is more expressive than the other model, based on a static attention



Attention in standard GAT



Attention in GATv2

- GATv2 is available as part of PyTorch Geometric library

```
from torch_geometric.nn.conv.gatv2_conv import GATv2Conv
```

[🏠](#) / [torch_geometric.nn](#) / [conv.GATv2Conv](#)

conv.GATv2Conv

```
class GATv2Conv ( in_channels: Union[int, Tuple[int, int]], out_channels: int, heads: int = 1, concat: bool = True, negative_slope: float = 0.2, dropout: float = 0.0, add_self_loops: bool = True, edge_dim: Optional[int] = None, fill_value: Union[float, Tensor, str] = 'mean', bias: bool = True, share_weights: bool = False, **kwargs ) \[source\]
```

Bases: `MessagePassing`

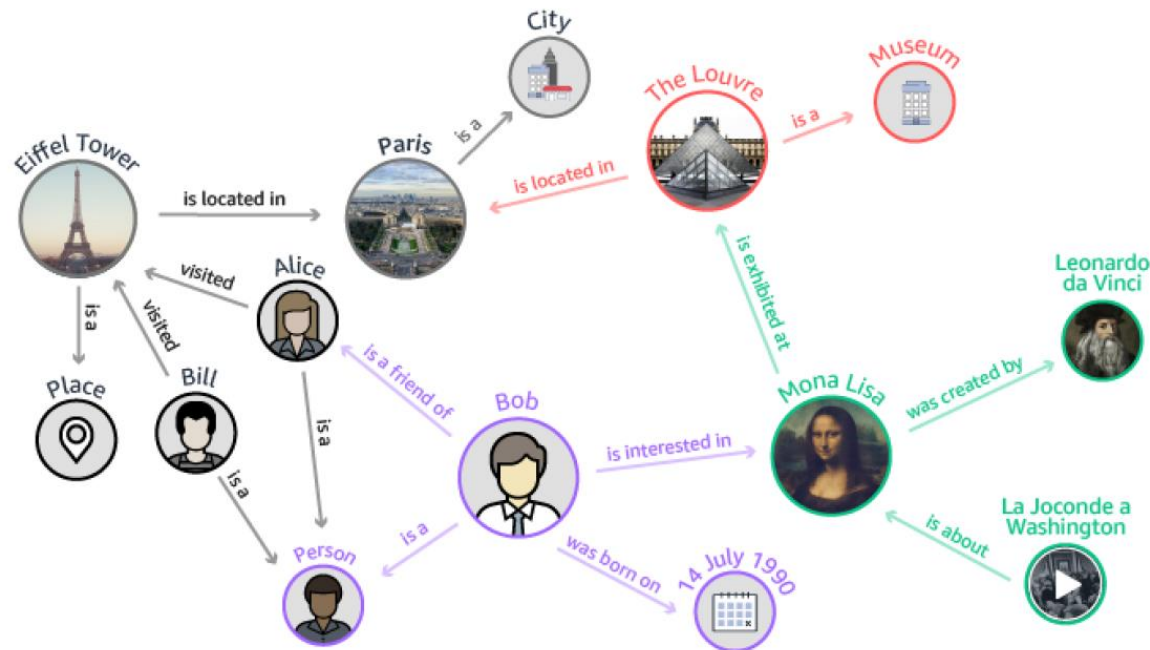
The GATv2 operator from the “How Attentive are Graph Attention Networks?” paper, which fixes the static attention problem of the standard `GATConv` layer. Since the linear layers in the standard GAT are applied right after each other, the ranking of attended nodes is unconditioned on the query node. In contrast, in `GATv2`, every node can attend to any other node.

- GATv2 is available as part of PyTorch Geometric library


```
from torch_geometric.nn.conv.gatv2_conv import GATv2Conv
```

```
136 class GATv2(torch.nn.Module):
137     def __init__(self, num_features, num_classes, dims, drop=0.0):
138         super(GATv2, self).__init__()
139         heads = 8
140         self.h = None
141         self.conv1 = GATv2Conv(num_features, dims, heads=heads, dropout = 0.3, concat=False)
142         self.conv2 = GATv2Conv(dims, num_classes, heads=heads, concat=False, dropout=0.3)
143         self.drop = torch.nn.Dropout(p=drop)
144     def forward(self, x, edge_index, g, Kindices):
145         x = F.elu(self.conv1(x, edge_index))
146         x = self.drop(x)
147         x = self.conv2(x, edge_index)
148         self.h = x
149         return F.log_softmax(x, dim=1)
```


- Graph in real world:
 - Many node types, link types
 - Non- ordered
 - Complex connections



- Multiple types of nodes or links
- Rich semantic information
 - Meta-path: a relation sequence connecting objects (e.g., Movie-Actor-Movie).



Homogeneous Graph



Heterogeneous Graph



Movie-Director-Movie

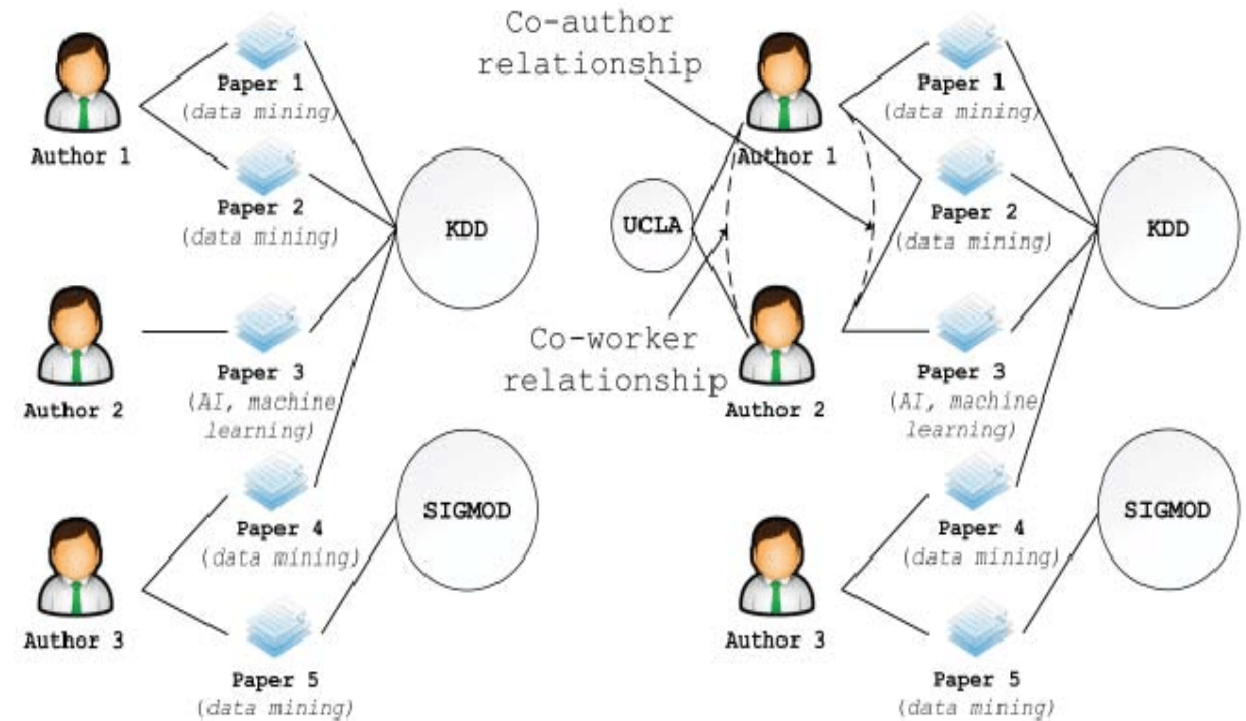
Two movies directed by the same director.



Movie-Actor-Movie

Two movies are starred by the same actor.

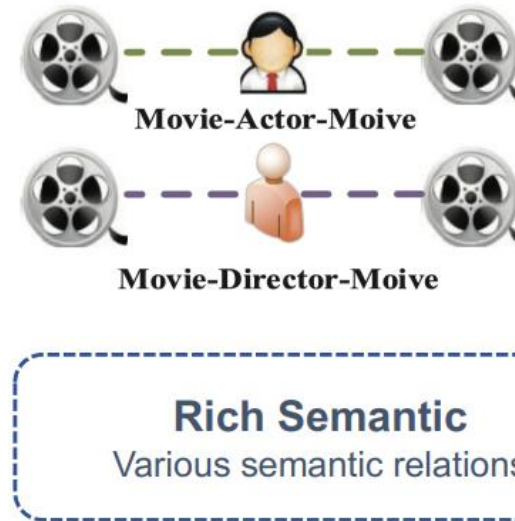

- DBLP Bibliographic network
 - Node (type)
 - KDD (Venue)
 - Author 1
 - Link (Type)
 - Write (Author - Paper)
 - Publish (Paper – Venue)



A. Examples of A-P-V-P-A meta-path on DBLP

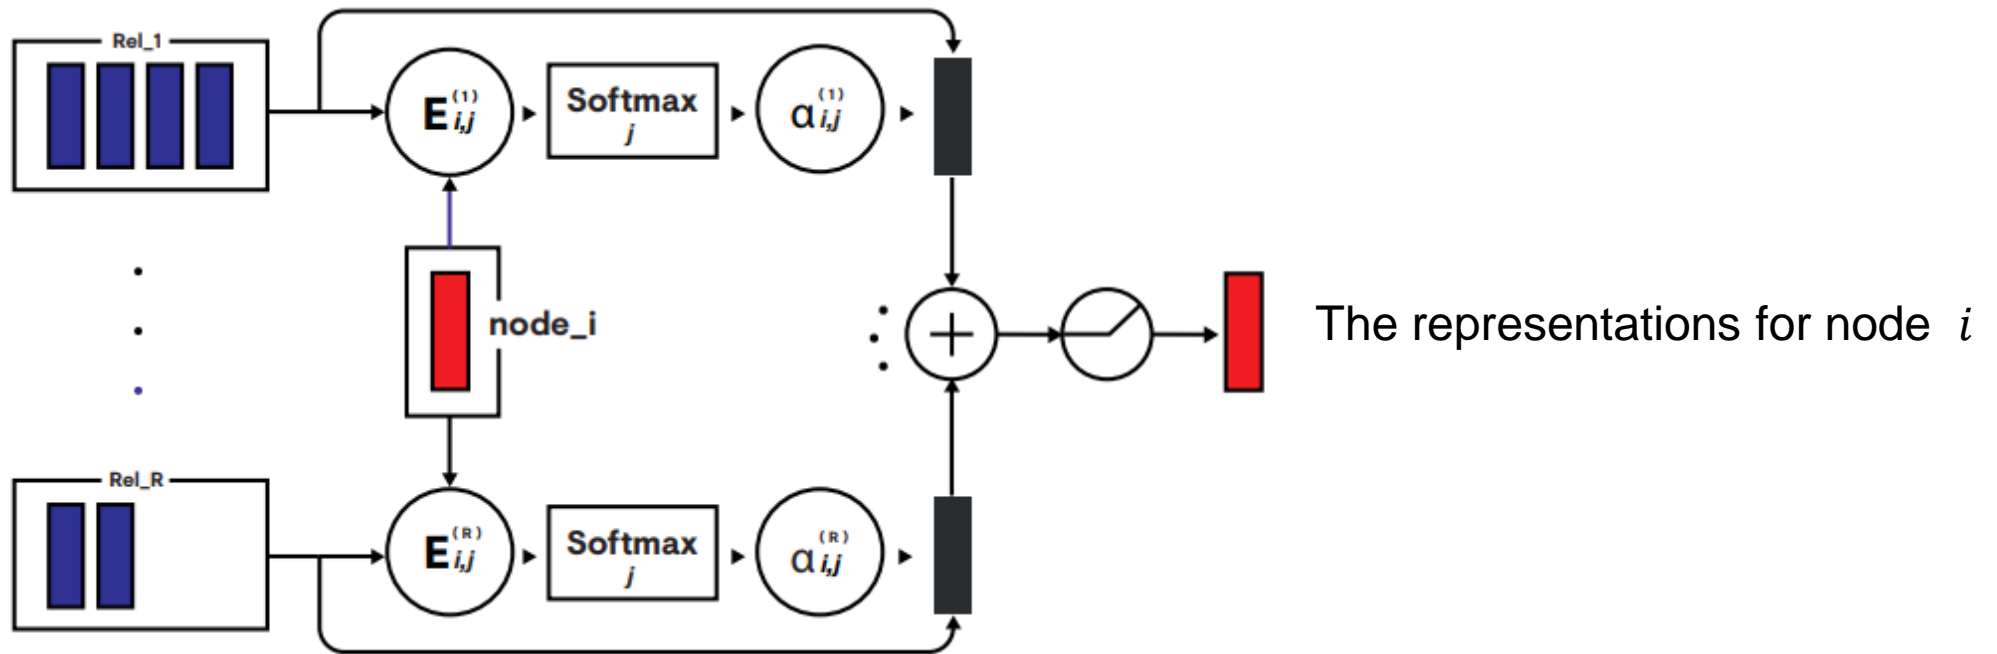
B. Examples of common neighborhood objects between two authors in DBLP

- Existing GNNs focus on homogeneous graphs
 - Cannot handle multiple types of nodes and edges.
 - Cannot capture rich semantic information.



Challenge: How to handle the heterogeneity of graph?

- The objective: Extending attention mechanisms to the relational graph domain



A target node i have different relations : $Rel_1, Rel_2, \dots, Rel_R$

The logits $E_{i,j}^{(r)}$ of each relation r : $E_{i,j}^{(r)} = a \left(g_i^{(r)}, g_j^{(r)} \right)$,

where: $G^{(r)} = H W^{(r)} \in \mathbb{R}^{N \times F'}$, the representation feature matrix under relation r

- GATv2 is available as part of PyTorch Geometric library

```
from torch_geometric.nn import RGATConv
```

```
16  ✓ class RGAT(torch.nn.Module):
17  ✓     def __init__(self, in_channels, hidden_channels, out_channels,
18             num_relations):
19             super().__init__()
20             self.conv1 = RGATConv(in_channels, hidden_channels, num_relations)
21             self.conv2 = RGATConv(hidden_channels, hidden_channels, num_relations)
22             self.lin = torch.nn.Linear(hidden_channels, out_channels)
23
24  ✓     def forward(self, x, edge_index, edge_type):
25             x = self.conv1(x, edge_index, edge_type).relu()
26             x = self.conv2(x, edge_index, edge_type).relu()
27             x = self.lin(x)
28             return F.log_softmax(x, dim=-1)
```



네트워크 과학연구실
NETWORK SCIENCE LAB



가톨릭대학교
THE CATHOLIC UNIVERSITY OF KOREA

